

C++ Expert training: C++ coroutines [**n4736**]

ir J.P. Lammertink

2020, March 4th

Program

- ▶ 13:00h Theory of coroutines
- ▶ 14:45h [snack]
- ▶ 15:00h Research workshop part 1
- ▶ 17:30h [dinner]
- ▶ 18:30h Research workshop part 2
- ▶ 21:00h End

Each research workshop part

- ▶ Choose and share your subject of interest
- ▶ Start individually or form a small team
- ▶ Research
- ▶ A round of 5 minutes presentation (individually)

Fixed presentation format

- ▶ My research subject is ...
- ▶ I discovered ...
- ▶ I still would like to know ...

Goals

- ▶ Learn about C++ coroutines [n4736]
- ▶ Learn about the CppCoro library
- ▶ Learn from each other
- ▶ Bonding
- ▶ Inspire
- ▶ Enjoy

My motivation

I'm wondering...

- ▶ frameworks in C++, such as COM, CORBA, SSCF, DDF, QT, ASIO
- ▶ they all impose a specific execution architecture
- ▶ ...is that OO?

My motivation

Theorem

A framework that includes:

- ▶ *serialization/deserializaion*
- ▶ *marshalling*
- ▶ *inter-process and inter-node communication*
- ▶ *asynchronous/overlapped IO/communication*
- ▶ *event processing*
- ▶ *TCP/IP and serial communication*

always consists of tightly coupled components.

Motivation

Theorem

C++20 has enough expression power to decompose such frameworks into loosely coupled components, without introducing (needless) overhead.

Sources

Lewis Baker

- ▶ 3 articles on GitHub [**Baker2018**]
- ▶ cppcoro library [**CppCoro**]

C++ Coroutines TS (N4736)

What is a coroutine

Definition

A coroutine is a generalisation of a function that allows the function to be *suspended* and then later *resumed*.

Definition

In C++20, a function is a coroutine, when it has at least one of the following keywords in its body:

- ▶ `co_await`
- ▶ `co_yield`
- ▶ `co_return`

From compiler-implementation-perspective, a function is a coroutine, for which the ‘activation frame’ solely resides on the stack.

‘Normal’ function call

<Call>

- ▶ Prepare caller for suspension
- ▶ Allocate memory for the activation frame
- ▶ Pass parameters
- ▶ Write address of resume-point of the caller
- ▶ Jump to the begin address of the function
- ▶ (Optionally, create local variables)

<Return>

- ▶ Optionally, store the return value
- ▶ Destroy local variables
- ▶ Optionally, destroy parameter objects
- ▶ Free memory for the coroutine activation frame
- ▶ Restore activation frame of caller
- ▶ Jump to the resume-point of caller

Stack



Stack

- ▶ A brief history of the stack; Sten Henriksson [**Henriksson**]
- ▶ Etymologically related haystack
- ▶ First design with stack: PERM II computer 1957
- ▶ First realisation: KDF9 computer 1959
- ▶ The push down automaton was introduced by Newell et al in 1959
- ▶ 'Stack' appears in first publication of E.W. Dijkstra in 1960 (before named 'pushdown store', 'LIFO list' or 'cellar storage')
- ▶ VAX computer implemented the CISC instructions pop and push 1978
- ▶ Enabler for recursive function calling

Example recursive Fibonacci

```
int recursive_fibonacci(int n)
{
    if (n <= 1) return n;
    else
    {
        int n_2 = recursive_fibonacci(n - 2);
        int n_1 = recursive_fibonacci(n - 1);
        return n_1 + n_2;
    }
}
```

Example recursive Fibonacci main

```
bool recursive_fibonacci_max(int n, int const max, int& f)
{
    if (n <= 1)
    {
        f = n;
        return f <= max;
    }
    else
    {
        int n_2 = recursive_fibonacci(n - 2);
        int n_1 = recursive_fibonacci(n - 1);
        f = n_1 + n_2;
        return n_2 < max - n_1;
    }
}

int main()
{
    for (int f, n = 0; recursive_fibonacci_max(n, 10'000, f); ++n)
    {
        std::cout << f << std::endl;
    }
}
```

Coroutine frame

The complete activation frame of a coroutine consists of two parts:

- ▶ Stack part
- ▶ Persistent part (often heap; accessible with handle)

<Suspend>

A coroutine function can be suspended in the middle of a function at suspension points indicated with `co_await` or `co_yield`

- ▶ Prepare coroutine for suspension
- ▶ Write address of resume-point
- ▶ (some programmer defined logic can be executed. \Rightarrow immediately resume/continue)
- ▶ The stack part of the coroutine activation frame is freed
- ▶ Restore activation frame of caller
- ▶ Jump to the resume-point of caller

<Resume>

A coroutine can be resumed by calling the `void resume()` method on the coroutine-frame handle.

- ▶ Prepare caller for suspension
- ▶ Allocate memory for stack part of the activation frame
- ▶ Write address of resume-point of the caller
- ▶ Jump to the resume-point of the coroutine

<Destroy>

The ‘destroy’ operation destroys the coroutine frame without resuming execution of the coroutine.

- ▶ Prepare caller for suspension
- ▶ Allocate memory for stack part of the activation frame
- ▶ Write address of resume-point of the caller
- ▶ Destruct all local variables in-scope
- ▶ Free memory for the coroutine activation frame (all)
- ▶ Restore activation frame of caller
- ▶ Jump to the resume-point of caller

Configuration: All Configurations ▾

Platform: All Platforms ▾

Configuration Manager...

▲ Configuration Properties

- General
- Advanced
- Debugging
- VC++ Directories
- ▲ C/C++
 - General
 - Optimization
 - Preprocessor
 - Code Generation
 - Language
 - Precompiled Headers
 - Output Files
 - Browse Information
 - Advanced
 - All Options
 - Command Line

- ▷ Linker
- ▷ Manifest Tool
- ▷ XML Document Generator
- ▷ Browse Information
- ▷ Build Events
- ▷ Custom Build Step
- ▷ Code Analysis

All Options

<different options>

Additional Options

Inherit from parent or project defaults ☒

/await

Example main

```
#include <iostream>

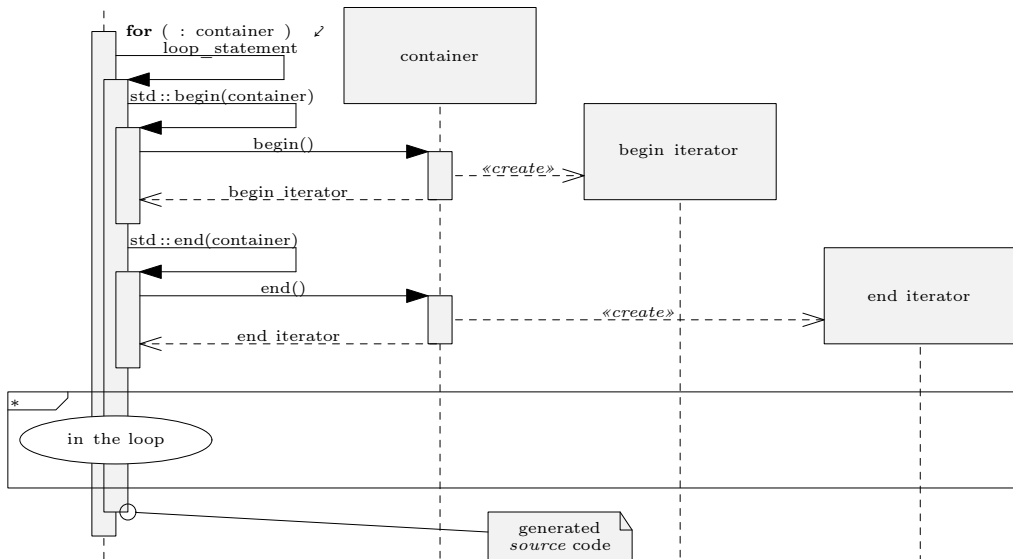
int main()
{
    int_generator_t int_generator(my_fibonacci_generator(10'000));
    int f;
    while (int_generator.get(f))
    {
        std::cout << f << std::endl;
    }
}
```

Example coroutine

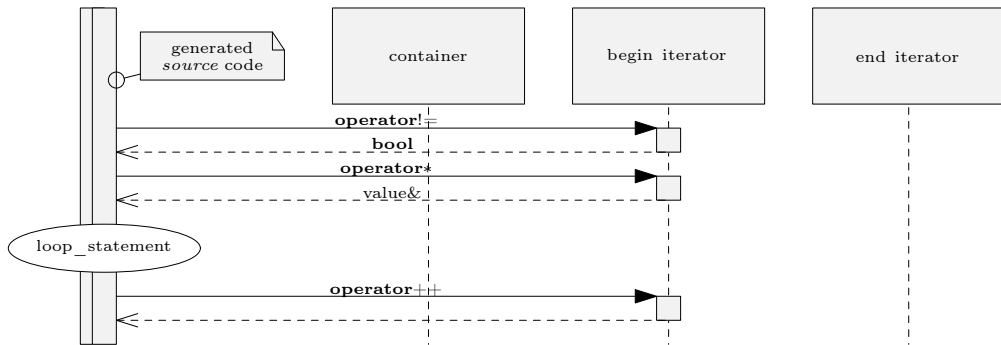
```
#include <experimental/resumable>

int_generator_t my_fibonacci_generator(int max)
{
    int n_2 = 0;
    if (max < n_2) co_return;
    co_yield n_2;
    int n_1 = 1;
    if (max < n_1) co_return;
    co_yield n_1;
    while (true)
    {
        if (max - n_1 < n_2) co_return;
        int n_0 = n_2 + n_1;
        co_yield n_0;
        n_2 = n_1;
        n_1 = n_0;
    }
}
```

Range-based for loop when 'range expression' is a container



Range-based for loop; in the loop



Range-based for loop; in code

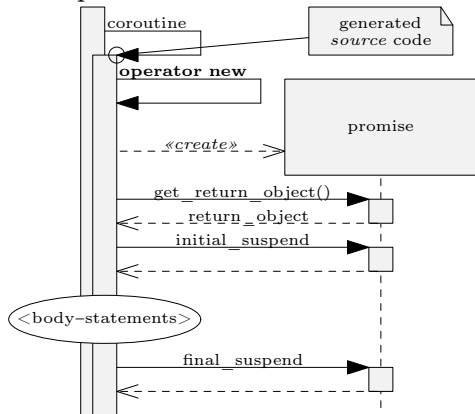
```
for ( init-statement(optional) range_declaration : range_expression )  
loop_statement
```

⇒

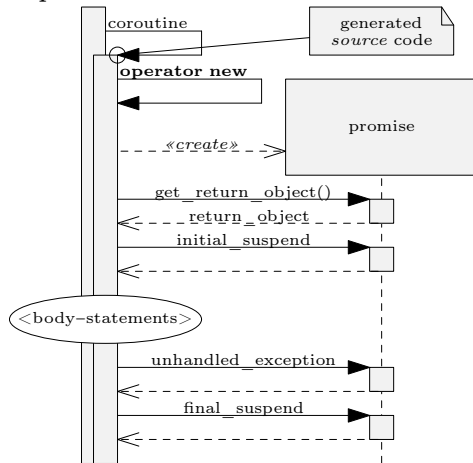
```
{  
    init-statement  
    auto && __range = range_expression ;  
    auto __begin = begin_expr ;  
    auto __end = end_expr ;  
    for ( ; __begin != __end; ++__begin)  
    {  
        range_declaration = *__begin;  
        loop_statement  
    }  
}
```

'Coroutine' function call

No exception



Exception



‘Coroutine’ function call

```
T some_coroutine(P param)
{
    auto* f = new coroutine_frame(std::forward<P>(param));
    auto returnObject = f->promise.get_return_object();
    // This call will return when the coroutine gets to the first
    // suspend-point or when the coroutine runs to completion.
    coroutine_handle<...>::from_promise(f->promise).resume();
    return returnObject;
}

coroutine_handle<...>::resume()
{
    co_await promise->initial_suspend();
    try
    {
        <body-statements>
    }
    catch (...)
    {
        promise->unhandled_exception();
    }
    FinalSuspend:
    co_await promise.final_suspend();
}
```

promise_type class

promise
<pre>(operator new(size_t)) ... get_return_object() void return_void() void return_value(... value) <awaiter> initial_suspend() (<awaitable> await_transform(... value)) (<awaiter> yield_value(... value)) (unhandled_exception()) <awaiter> final_suspend()</pre>

Notion of “promise”

- ▶ Sometimes “promise” is like `std::promise` part of promise-future pair
reminder:
calling `set_value(...)` on `std::promise`
results in `get()` to return on `std::future`
- ▶ For some use cases “promise” is *not* like `std::promise`
- ▶ \Rightarrow “promise” is like “coroutine state controller”

Example `promise_type`

```
struct promise_type
{
    int_generator_t get_return_object()
    { ... }
    ... initial_suspend()
    { ... }
    ... final_suspend()
    { ... }
    ...
};
```

Two options to obtain the `promise_type`

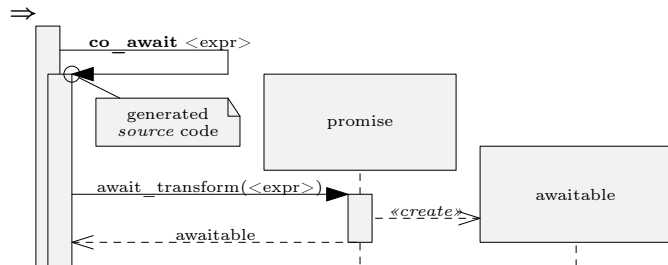
1. Return type of coroutine/function has member-type `promise_type`
2. Using `coroutine_traits`
assume coroutine `foo` with signature: `T0 foo(T1, T2, ..., Tn)`
then `promise_type` can be defined under:
`typename coroutine_traits<T0, T1, T2, ..., Tn>::` ↗
`promise_type;`

Example promise_type class

```
struct int_generator_t
{
    struct promise_type
    {
        using coro_handle_t =
            std::experimental::coroutine_handle<promise_type>;
        int_generator_t get_return_object()
        {
            return int_generator_t(coro_handle_t::from_promise(*this));
        }
        ...
    };
    int_generator_t(promise_type::coro_handle_t coro_handle)
        : m_coro_handle(coro_handle)
    {}
    promise_type::coro_handle_t m_coro_handle;
};
```


`co_await <expr>`

`await_transform` exists



`await_transform` does
not exist

⇒
`awaitable = <expr>`

`co_await <expr>`

`co_await <expr>`


\Rightarrow

```
template<typename P, typename T>
decltype(auto) get_awaitable(P& promise, T&& expr)
{
    if constexpr (has_any_await_transform_member_v<P>)
        return promise.await_transform(expr);
    else
        return expr;
}
...

{
    auto&& value = <expr>;
    auto&& awaitable = get_awaitable(promise, value);
    ...
}
```

operator `co_await` awaitable

operator `co_await` exists



operator `co_await` does *not* exist

⇒

`awaiter =` ↴
`awaitable`

operator co_await awaitable

operator co_await


⇒

```
template<typename Awaitable>
decltype(auto) getawaiter(Awaitable&& awaitable)
{
    if constexpr (has_member_operator_co_await_v<Awaitable>)
        return awaitable.operator co_await();
    else if constexpr (has_non_member_operator_co_await_v<Awaitable  ↵
        &&>)
        return operator co_await(awaitable);
    else
        return awaitable;
}

...

{
    auto&& awaitable = getawaitable(promise, value);
    auto&& awaiter = getawaiter(awaitable);
```

await_ready returning true

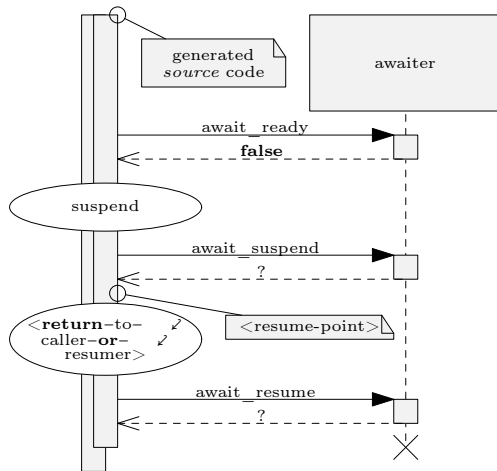


await_ready returning true

```
{  
    auto&& value = <expr>;  
    auto&& awaitable = get_awaitable(promise, value);  
    auto&& awaiter = get_awaiter(awaitable);  
    if (!awaiter.await_ready())  
    {  
        ... /* skipped when await_ready() returns true */  
    }  
    return awaiter.await_resume();  
}
```

Note that `co_await <expr>` may return a value, when `await_resume()` does.

await_ready returning false



await_ready returning false

```
{
    ...
    if (!awaiter.await_ready())
    {
        using handle_t = std::experimental::coroutine_handle<P>;
        using await_suspend_result_t = decltype(awaiter.await_suspend( ↵
            handle_t::from_promise(p)));
        <suspend-coroutine>
        if constexpr (std::is_void_v<await_suspend_result_t>)
        {
            awaiter.await_suspend(handle_t::from_promise(p));
            <return-to-caller-or-resumer>
        }
        else
        {
            ...
        }
    }
    <resume-point>
}
return awaiter.await_resume();
}
```


await_ready returning false, await_suspend returning bool

```
{
    ...
    if (!awaiter.await_ready())
    {
        using handle_t = std::experimental::coroutine_handle<P>;
        using await_suspend_result_t = decltype(awaiter.await_suspend( ↵
            handle_t::from_promise(p)));
        <suspend-coroutine>
        if constexpr (std::is_void_v<await_suspend_result_t>)
            ...
        else
        {
            if (awaiter.await_suspend(handle_t::from_promise(p)))
            {
                <return-to-caller-or-resumer>
            }
        }
        <resume-point>
    }
    return awaiter.await_resume();
}
```

awaiter class

awaiter
<pre>bool await_ready() ... await_suspend(coroutine_handle<>) ... await_resume()</pre>

Example usage of awaiter class

```
struct int_generator_t
{
    struct promise_type
    {
        std::experimental::suspend_always initial_suspend()
        {
            return std::experimental::suspend_always();
        }
        std::experimental::suspend_always final_suspend()
        {
            return std::experimental::suspend_always();
        }
        ...
    };
    ...
};
```

std::experimental::suspend_always

```
struct suspend_always
{
    bool await_ready() noexcept
    {
        return false;
    }
    void await_suspend(coroutine_handle<>) noexcept {}
    void await_resume() noexcept {}
};
```

co_yield

co_yield <expr>


translates into

co_await promise.yield_value(<expr>)


Example co_yield

```
struct int_generator_t
{
    struct promise_type
    {
        ...
        std::experimental::suspend_always yield_value(int v)
        {
            m_v = v;
            return std::experimental::suspend_always();
        }
        ...
        int m_v;
    };
    ...
};
```

`co_return`



`goto FinalSuspend;`



`goto FinalSuspend;`

If execution runs off the end of a coroutine without a `co_return` statement, then this is equivalent to having a `co_return;` at the end of the function body.

Example return_void

```
struct int_generator_t
{
    struct promise_type
    {
        ...
        void return_void() {}
        ...
    };
    ...
};
```


Handle and promise

- ▶ Handle \Rightarrow promise: `m_coro_handle.promise()`
- ▶ Promise \Rightarrow handle: `coro_handle_t::from_promise(*this)`
- ▶ Resuming coroutine: `coro_handle.resume()`
- ▶ Check wheter coroutine finished: `coro_handle.done()`

Example handle and promise

```
struct int_generator_t
{
    struct promise_type
    {
        using coro_handle_t = std::experimental::coroutine_handle< ↗
            promise_type>;
        int_generator_t get_return_object()
        {
            return int_generator_t(coro_handle_t::from_promise(*this));
        }
        ...
        int m_v;
    };
    ...
    int_generator_t(promise_type::coro_handle_t coro_handle)
        : m_coro_handle(coro_handle)
    {}
    bool get(int& v)
    {
        m_coro_handle.resume();
        v = m_coro_handle.promise().m_v;
        return !m_coro_handle.done();
    }
    promise_type::coro_handle_t m_coro_handle;
};
```

Example output

0
1
1
2
3
5
8
13
21
34
55
89
144
233
377
610
987
1597
2584
4181
6765

Example fancy main

```
int main()
{
    for (auto f : my_fibonacci_generator(10'000))
    {
        std::cout << f << std::endl;
    }
}
```

Example range based for logic

```
struct int_generator_t
{
    struct promise_type
    {
        ...
        bool resume()
        {
            coro_handle_t coro_handle(coro_handle_t::from_promise(* ↵
                this));
            coro_handle.resume();
            return !coro_handle.done();
        }
    };
struct iterator
{
    iterator(promise_type& promise) : m_promise(promise) {}
    bool operator!=(int) const { return m_promise.resume(); }
    int operator*() const { return m_promise.m_v; }
    iterator operator++() { return *this; }
    promise_type& m_promise;
};
...
iterator begin() { return iterator(m_coro_handle.promise()); }
int end() { return 0; }
};
```

cppcoro::generator

```
#include <cppcoro/generator.hpp>
```

```
cppcoro::generator<int> fibonacci_generator(int max)
```

```
{  
    int n_2 = 0;  
    if (max < n_2) co_return;  
    co_yield n_2;  
    int n_1 = 1;  
    if (max < n_1) co_return;  
    co_yield n_1;  
    while (true)  
    {  
        if (max - n_1 < n_2) co_return;  
        int n_0 = n_2 + n_1;  
        co_yield n_0;  
        n_2 = n_1;  
        n_1 = n_0;  
    }  
}
```

```
int main()
```

```
{  
    for (auto f : fibonacci_generator(10'000))  
    {  
        std::cout << f << std::endl;  
    }  
}
```

cppcoro

► Coroutine Types

- `task<T>`
- `shared_task<T>`
- *generator*`<T>`
- `recursive_generator<T>`
- `async_generator<T>`

► Awaitable Types

- `single_consumer_event`
- `single_consumer_async_auto_reset_event`
- `async_mutex`
- `async_manual_reset_event`
- `async_auto_reset_event`
- `async_latch`
- `sequence_barrier`
- `multi_producer_sequencer`
- `single_producer_sequencer`

cppcoro

- ▶ Functions
 - ▶ `sync_wait()`
 - ▶ `when_all()`
 - ▶ `when_all_ready()`
 - ▶ `fmap()`
 - ▶ `schedule_on()`
 - ▶ `resume_on()`
- ▶ Cancellation
 - ▶ `cancellation_token`
 - ▶ `cancellation_source`
 - ▶ `cancellation_registration`
- ▶ Schedulers and I/O
 - ▶ `static_thread_pool`
 - ▶ `io_service` and `io_work_scope`
 - ▶ `file`, `readable_file`, `writable_file`
 - ▶ `read_only_file`, `write_only_file`, `read_write_file`
- ▶ Networking
 - ▶ `socket`
 - ▶ `ip_address`, `ipv4_address`, `ipv6_address`
 - ▶ `ip_endpoint`, `ipv4_endpoint`, `ipv6_endpoint`

Speed test

- ▶ 5 pieces of work: 1, 3, 5 interleave with 2 and 4
- ▶ implementation with hardcoded interleaving
- ▶ implementation with coroutines and events
- ▶ implementation with `std::async`

Speed test output

test 1, hardcoded interleaving:	600ns
test 2, coroutine event implementation:	1100ns
test 3, async implementation:	12000ns
async to coroutine ratio:	22.8

With custom new operator

test 1, hardcoded interleaving:	600ns
test 2, coroutine event implementation:	1000ns
test 3, async implementation:	12000ns
async to coroutine ratio:	28.5

References

- [1] [Baker, 2018] Lewis S. Baker.
Asymmetric Transfer
<https://lewissbaker.github.io/2017/09/25/coroutine-theory>
<https://lewissbaker.github.io/2017/11/17/understanding-operators>
<https://lewissbaker.github.io/2018/09/05/understanding-the-problem>
- [2] [CppCoro] Lewis S. Baker.
CppCoro - A coroutine library for C++
- [3] [n4736]
<http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2018/n4736>
- [4] [Henriksson] Sten Henriksson.
A brief history of the stack
Computer Science Department, Lund University, Lund, Sweden

Tech talks

- ▶ CppCon 2016: James McNellis “Introduction to C++ Coroutines”
<https://www.youtube.com/watch?v=ZTqHjjm86Bw>
- ▶ CppCon 2016: Gor Nishanov “C++ Coroutines: Under the covers”
<https://www.youtube.com/watch?v=8C8NnE1Dg4A>